

Sharpe-Ratio Inference at the Date Boundary in Financial Panels

Abstract

Using public Kenneth French panels, AQR factor series, and Monte Carlo designs, this paper shows that Sharpe-ratio evidence can be overstated when selected asset-date rows are treated as independent even though the payoff is a date-level portfolio return. I compare row-pooled IID Sharpe uncertainty with date-level long-run Sharpe uncertainty through Sharpe UVIF and evaluate claims using HAC-delta and dependent-bootstrap inference on date returns. In dependent-null simulations, row-naive tests reject about 35 percent of the time at a nominal 5 percent level, while date-level procedures remain close to nominal size. The public benchmarks illustrate when the correction preserves strong momentum evidence and when weaker panel claims disappear.

Keywords: Sharpe ratio; panel data; cross-sectional dependence; serial dependence; HAC/bootstrap inference; portfolio evaluation

JEL Codes: C12; C15; G11; G12; G17

1 Introduction

1.1 Problem: Row-to-date sampling boundary

Empirical finance often evaluates signals in entity-time panels: factor screens, anomaly sorts, automated portfolio-signal screens, model-selection grids, and trading-strategy backtests. The practical question in this paper is whether the reported Sharpe evidence is being attached to the right sampling unit. In the Monte Carlo null designs reported below, treating selected rows as independent produces rejection rates around 35 percent at a nominal 5 percent level when the true edge is zero, while date-level HAC and dependent-bootstrap methods stay much closer to nominal size.

Suppose a signal screen selects many assets on each decision date and then reports a Sharpe ratio by pooling all selected asset-date rows. If the payoff is the portfolio return earned on each date, the apparent row count can exaggerate precision even before

serial dependence is considered. A row-pooled Sharpe report or t -test implicitly treats the cross-section on a date as many independent draws from the same experiment. A portfolio signal or trading strategy, however, is marked to market by date.

I denote the underlying entity-time panel by $\mathcal{P} = \{(t, i, s_{t,i}, r_{t+1,i})\}$, where t indexes dates, i indexes assets or portfolios, $s_{t,i}$ is the signal, and $r_{t+1,i}$ is the next-period return. One evaluation failure mode is to evaluate a positive edge using all selected asset-date rows as though they were independent. This turns a panel with T dates and n_t selected rows per date into a nominal sample of $\sum_t n_t$, even when same-date rows share common shocks. The row-pooled construction is retained only as a diagnostic baseline against which date-level inference is compared.

The issue extends beyond a single trading-strategy application. Factor screens, anomaly panels, model-selection procedures, and automated portfolio-signal evaluations can all multiply observations across assets on the same date while the economic payoff is realized through a date-level portfolio return. The paper uses public panel and factor examples because they make the sampling boundary visible, but the statistical problem is broader: what uncertainty should be attached to a Sharpe estimate after cross-sectional row multiplication and serial dependence in the date-return series are recognized. Cross-sectional dependence determines how much extra information the additional selected rows actually contain, and serial dependence determines the long-run variance of the date series. Standard clustering techniques, including the finance-panel guidance of Petersen (2009), address uncertainty for linear means or regression coefficients; they do not by themselves map row-pooled Sharpe evidence to the long-run variance of the economically relevant date-level Sharpe functional.

Automated empirical signal screens can evaluate large numbers of entity-date rows and many menu variants in a single run. The distortion is largest when many same-date entities are selected, common shocks are strong, menus are searched, and the confirmatory statistic is a nonlinear Sharpe functional.

The intuition is simple. If 50 selected rows are generated from only a small number of economically relevant dates, the row count can make the Sharpe estimate look more precise than the portfolio-return evidence supports. The relevant sampling object is the date-level return series, not the multiplied row count. Sharpe UVIF records how far the row-pooled IID uncertainty sits from the date-level long-run uncertainty.

The required diagnostic is a ratio between the uncertainty attached to the date-level Sharpe claim and the uncertainty a row-pooled IID analysis would report. I refer to this ratio as the Sharpe Unified Variance Inflation Factor, or Sharpe UVIF:

$$\text{UVIF}_{\text{SR}} = \frac{\text{Var}_{LR}(\widehat{\text{SR}}_{\text{date}})}{\text{Var}_{\text{IID}, \text{row}}(\widehat{\text{SR}}_{\text{row}})}.$$

It is best interpreted as a sampling-boundary distortion ratio for reported Sharpe uncer-

tainty. The numerator is the uncertainty of the economically relevant date-level Sharpe claim. The denominator is the uncertainty that would be reported by a row-pooled IID analysis of the selected entity-date returns. When the row-pooled and date-level estimators target the same Sharpe functional under exchangeability and equal-weight aggregation, Sharpe UVIF reduces to a design-effect variance inflation factor for the non-linear Sharpe ratio. Sharpe UVIF is interpreted as a reported-uncertainty ratio between row-pooled IID Sharpe inference and date-level long-run Sharpe inference.

1.2 Contribution and empirical approach

This paper contributes to financial econometrics in three ways. First, using public Kenneth French panels, AQR factor series, and Monte Carlo designs, it shows that the correction preserves canonical momentum evidence while weaker stress-panel claims can disappear after the evidence is moved to date returns. Second, it shows that Sharpe-ratio evidence from selected panel rows can be distorted when same-date observations are treated as independent even though the relevant payoff is a date-level portfolio return. Third, it proposes Sharpe UVIF, a diagnostic that compares row-pooled reported Sharpe uncertainty with date-level long-run uncertainty and extends standard design-effect logic to the nonlinear Sharpe functional.

The empirical section checks the implementation on canonical French WML, evaluates French panel stress cases, and reports AQR factor benchmarks as time-series-only evidence. The public Kenneth French and AQR evidence is based on documented public data sources and reproduction files. The conservative empirical question is whether apparently significant portfolio-signal claims owe their precision to row-level sampling and disappear after date aggregation, dependence adjustment, and researcher-menu correction. The replication code also runs a Monte Carlo study comparing row-naive testing with date-IID bootstrap, HAC-delta, moving-block bootstrap, stationary bootstrap, and joint resampling decisions.

2 Literature and motivation

This paper sits at the intersection of literatures that are often discussed separately. Survey and grouped-error design effects explain why row counts overstate information (Kish, 1965; Moulton, 1986), and finance-panel standard error work explains why cross-sectional and time clustering matter in empirical finance (Petersen, 2009). Those results are central background, but they mainly address means or regression coefficients rather than nonlinear portfolio-performance evidence after entity-date rows have been aggregated into a date-level return. They do not define a target-aware distortion ratio for nonlinear portfolio performance statistics.

Sharpe-ratio inference follows a separate time-series literature. Lo (2002) studies Sharpe-ratio sampling under return dynamics, and robust Sharpe testing builds on HAC and bootstrap methods (Ledoit and Wolf, 2008; Opdyke, 2007; Andrews, 1991; Kiefer and Vogelsang, 2005; Kunsch, 1989; Politis and Romano, 1994; Lahiri, 2003). These tools define the date-level sampling law once the return series is specified. They do not quantify how much apparent Sharpe precision was created before the evidence was moved from selected entity-date rows to the economically relevant date-return boundary. The boundary problem occurs before Sharpe inference: which return series is the object? What is missing in the Sharpe-inference literature is an explicit mapping from row-pooled Sharpe evidence to date-level long-run Sharpe evidence.

Researcher-menu and factor-zoo work addresses a third source of false precision. White Reality Check and Romano-Wolf procedures control data-snooping and familywise error in menus (White, 2000; Romano and Wolf, 2005), while the empirical asset-pricing literature documents false discoveries, factor-alpha fragility, anomaly decay, and trading costs (Harvey et al., 2016; Harvey and Liu, 2020a,b; Bailey and Lopez de Prado, 2014; Holm, 1979; Fama and French, 1993; Carhart, 1997; Fama and French, 2010; McLean and Pontiff, 2016; Almgren and Chriss, 2001; Hasbrouck, 2009; Frazzini et al., 2015). This paper is adjacent to that literature but distinct: its object is the sampling unit and return boundary before a row-level Sharpe claim is interpreted as portfolio-performance evidence. It separates three failure channels that are often bundled together in practice: cross-sectional row multiplication, serial dependence in the date-return series, and researcher-menu search.

Recent empirical asset-pricing work asks whether published predictors survive stricter implementation, data, and multiple-testing standards. The present paper asks an upstream question: before a predictor is interpreted as portfolio-performance evidence, has the sampling unit been moved from selected entity-date rows to the date-level return that can be interpreted as portfolio-performance evidence?

3 Data and empirical design

The empirical design uses public Kenneth French momentum and Size/book-to-market portfolio panels as row-level applications, and public AQR factor series as pre-aggregated time-series benchmarks. The French row-level panels allow same-date row aggregation, selected-count diagnostics, and same-date placebo checks. The AQR series are useful factor-return benchmarks, but they do not contain constituent signal-return rows and are therefore reported as time-series evidence only. These public applications are benchmark illustrations of the method, not a survey of all asset-pricing anomalies.

The method is a dependence-aware inference design rather than a single test. Its first channel is cross-sectional: all evidence is translated from selected asset-date rows into one

economically interpretable date-level portfolio return, so same-date clusters cannot create sample size by row multiplication. Its second channel is serial: all Sharpe decisions are made on the date-return series using HAC-delta and dependent-bootstrap procedures. For panel candidates with row-level signal-return pairs, the empirical implementation reports both channels before adding researcher-menu, placebo, factor-alpha, and turnover-cost sensitivity diagnostics.

Each panel candidate is reported at three levels: the row-naive baseline, the date-level confirmatory estimate, and the UVIF diagnostic gap between them.




Figure 1: Inference workflow. The empirical claim is evaluated at the date-level portfolio-return boundary; row-pooled evidence is retained as a comparator rather than as the maintained sampling law.

3.1 Date aggregation as the target boundary

Let S_t denote the selected entities on date t , let $Y_{t+1,i}$ be the signed one-period row return, and let $w_{t,i}$ be the economic weight assigned to selected row i . The portfolio return used for Sharpe inference is the date-level object

$$R_{t+1} = \sum_{i \in S_t} w_{t,i} Y_{t+1,i}, \quad \sum_{i \in S_t} |w_{t,i}| = 1 \text{ on active public-panel dates.} \quad (1)$$

This is the target boundary. A row-pooled analysis estimates uncertainty as though the selected $Y_{t+1,i}$ were independent observations, while the confirmatory analysis estimates uncertainty for the realized date-return series $\{R_t\}$. In the empirical implementation, weights are normalized to the declared gross-exposure convention before returns are computed. Unless otherwise stated, rendered public panel tables use unit gross exposure within the selected signed-return set. For signed long-short rows, the selected positive and negative signs are combined into one unit-gross portfolio on each active date; dates with no selected row receive no active portfolio return.

In the public equal-weight threshold implementation, optional non-negative weights $a_{t,i}$ are normalized within the selected set S_t , where S_t is produced by the declared signal

160 and threshold rule, giving

$$\bar{r}_{t+1}(L, q, \tau) = \frac{\sum_{i \in \mathcal{S}_t} a_{t,i} y_{t+1,i}^{(L,q)}}{\sum_{i \in \mathcal{S}_t} a_{t,i}}. \quad (2)$$

161 The denominator is the total active economic weight on date t , not a generic row count.
 162 Under the primary equal-weighted evaluation $a_{t,i} = 1$, it reduces to the number of selected
 163 rows. Confidence-weighted and signal-magnitude-weighted returns are sensitivity checks.

164 3.2 Sharpe inference

165 The inference design separates estimation from testing. It reports two-sided 95 percent
 166 intervals for $\widehat{\text{SR}}$ and one-sided positive-edge p -values for $H_0 : \text{SR} \leq 0$. For a fixed return
 167 series with positive volatility, this is equivalent to a nonpositive mean excess return null.
 168 Bootstrap inference resamples the date-return series by IID date, moving block, or sta-
 169 tionary bootstrap. The date-IID bootstrap is reported only as a comparator; dependence-
 170 aware conclusions rely on HAC, moving-block, and stationary-bootstrap procedures. The
 171 positive-edge bootstrap p -value uses one-draw smoothing:

$$p^+ = \frac{1 + \#\{\widehat{\text{SR}}_b^* \leq 0\}}{B + 1}. \quad (3)$$

172 For HAC inference, reindex the finite realized date-return observations as R_u , $u =$
 173 $1, \dots, T_R$, where R_u is the strategy excess return for the decision date after subtracting
 174 the applicable risk-free rate when the source is not already a long-short excess-return
 175 series. Let A_f denote the source-specific annualization factor: $A_f = 252$ for daily sources
 176 and $A_f = 12$ for monthly sources in the public configuration. Let $g_u = (R_u, R_u^2)'$, $\hat{\mu} =$
 177 $T_R^{-1} \sum_u R_u$, $\hat{m}_2 = T_R^{-1} \sum_u R_u^2$, and $\hat{v} = \hat{m}_2 - \hat{\mu}^2$. The denominator uses the T_R -moment
 178 variance, not the $T - 1$ bias correction, because it is the plug-in moment entering the
 179 smooth functional

$$f(\mu, m_2) = \sqrt{A_f} \frac{\mu}{(m_2 - \mu^2)^{1/2}}.$$

180 This convention only changes finite-sample scaling and is kept consistent across the plug-
 181 in variance calculations. Following Ledoit and Wolf (2008), the Sharpe ratio is treated
 182 as a smooth functional $f(\theta)$ of the moment vector $\theta = \mathbb{E}[g_t]$. On sets where $m_2 -$
 183 μ^2 is bounded away from zero, this map is Hadamard differentiable. The asymptotic
 184 normality of $\sqrt{T_R}(\widehat{\text{SR}} - \text{SR})$ follows from the functional delta method provided the long-
 185 run covariance matrix Ω is consistently estimated by a kernel-based HAC estimator.
 186 Evaluated at $(\hat{\mu}, \hat{m}_2)$,

$$\widehat{\text{SR}} = \sqrt{A_f} \frac{\hat{\mu}}{\sqrt{\hat{v}}}, \quad \nabla f(\hat{\mu}, \hat{m}_2) = \left(\sqrt{A_f} \frac{\hat{m}_2}{\hat{v}^{3/2}}, -\frac{1}{2} \sqrt{A_f} \frac{\hat{\mu}}{\hat{v}^{3/2}} \right)'.$$

187 With $\widehat{\Omega}$ the Bartlett-kernel HAC estimate of the long-run covariance of g_u , the HAC-delta
 188 variance is

$$\widehat{\mathbb{V}\text{ar}}(\widehat{\text{SR}}) = \nabla f(\hat{\mu}, \hat{m}_2)' \widehat{\Omega} \nabla f(\hat{\mu}, \hat{m}_2) / T_R. \quad (4)$$

189 Explicitly, for centered moments $\tilde{g}_u = g_u - \bar{g}$,

$$\widehat{\Gamma}_k = T_R^{-1} \sum_{u=k+1}^{T_R} \tilde{g}_u \tilde{g}'_{u-k}, \quad \widehat{\Omega} = \widehat{\Gamma}_0 + \sum_{k=1}^{K_T} \left(1 - \frac{k}{K_T + 1}\right) (\widehat{\Gamma}_k + \widehat{\Gamma}'_k).$$

190 This is a covariance matrix for the joint process (R_u, R_u^2) , so the off-diagonal terms
 191 carry the long-run covariance between the sample mean and second moment into the
 192 Sharpe-ratio delta method. The bandwidth K_T is a tuning choice and may depend on
 193 T . The default HAC bandwidth uses the Andrews automatic Bartlett rule, with fixed-
 194 lag and prewhitened sensitivities reported in the appendix. Important empirical claims
 195 should be checked against reported nearby bandwidths and against the block-bootstrap
 196 intervals rather than relying on a single HAC lag choice. The technical appendix reports
 197 prewhitened HAC and fixed-fraction HAC sensitivity checks.

198 Because financial date returns can be heavy-tailed, HAC-delta evidence is reported
 199 alongside bootstrap p -values and intervals. When HAC and dependent-bootstrap conclu-
 200 sions disagree, the result is reported as robustness-fragile rather than summarized as a
 201 positive-edge finding.

202 4 Implementation and decision rules

203 4.1 Public-panel momentum implementation

204 The following construction is used for the public momentum-panel implementation; the
 205 sampling-boundary diagnostic itself does not depend on this particular signal. For look-
 206 back L and skip q , the skipped compounded signal is

$$s_{t,i}^{(L,q)} = \prod_{j=t-L-q+1}^{t-q} (1 + r_{j,i}) - 1. \quad (5)$$

207 With $q \geq 1$, the signal excludes the outcome return and is executable at date t when all
 208 signals in the available universe A_t are known before the trading decision. The signed
 209 one-period row return is

$$y_{t+1,i}^{(L,q)} = \text{sign}(s_{t,i}^{(L,q)}) r_{t+1,i}, \quad (6)$$

with $\text{sign}(0) = 0$; zero-signal and missing-signal rows receive zero economic weight. The confidence score $c_{t,i}^{(L,q)}$ is the within-date percentile rank of $|s_{t,i}^{(L,q)}|$, and the selected set is

$$\mathcal{S}_t(L, q, \tau) = \{i : c_{t,i}^{(L,q)} \geq \tau\}.$$

The percentile score is a baseline screening rule, not an optimal weighting theorem. Tie handling, one-name dates, and selected-count granularity are reported in the appendix and replication files.

4.2 Researcher menus, placebos, and factor alpha

For a fixed menu $\mathcal{M}$ of lookbacks and thresholds, the primary joint adjustment is Romano-Wolf stepdown max- t resampling over $\{R_u(m) : m \in \mathcal{M}\}$. Holm-adjusted marginal bootstrap p -values and White’s Reality Check are secondary diagnostics. The Deflated Sharpe Ratio diagnostic of Bailey et al. (2014) is reported as an additional selection-bias check. In the public implementation, the menu is the finite pre-specified set

$$\mathcal{M} = \{(c, \theta_c, \tau) : c \in \mathcal{C}, \tau \in \mathcal{T}_c\},$$

where θ_c denotes the candidate’s fixed signal definition. For panel candidates, $\mathcal{T}_c = \{0.0, 0.3, 0.5, 0.7, 0.9\}$ and any lookback and skip values are fixed before the confirmatory run rather than selected after seeing results. Pre-aggregated AQR factor-series candidates use the singleton threshold menu $\mathcal{T}_c = \{0.0\}$; because they do not contain row-level constituent signals, same-date panel permutation is not applicable to them. Real researcher menus are usually correlated rather than independent, so the procedure uses Romano-Wolf stepdown to control familywise error while preserving the joint dependence among menu elements.

The fixed-menu condition is a confirmatory convention, not a claim that research is never iterative. In practice, exploratory work should be separated from the final evaluation by freezing the signal definitions, lookbacks, thresholds, filters, cost assumptions, and benchmark model before the confirmatory run. If discarded variants influenced the final choice, they belong in the researcher menu or the exercise should be reported as exploratory. A stronger design uses a locked holdout period or an independently refreshed sample after the menu is frozen.

The public placebo is a same-date permutation design. For each date, the vector of signals is randomly permuted across portfolios, preserving the date’s cross-sectional signal distribution and missingness while breaking the signal-return pairing. The signed next-day returns, confidence ranks, selected counts, turnover, and date aggregation are then recomputed. Repeating this procedure gives a permutation null for the observed public momentum panel.

This placebo is a test under a within-date exchangeability hypothesis: under the null, the assignment of signals $s_{t,i}$ to next-period returns $r_{t+1,i}$ is exchangeable within each date t . If signals are structurally tied to size, sector, liquidity, country, exchange, beta, or other common-factor state variables, an unrestricted same-date permutation can break the wrong null. In those cases the analysis should permute within pre-specified blocks, residualize returns against the structural variables before permutation, or mark the placebo as not applicable. The permutation check is therefore a design-specific diagnostic rather than a universal validity guarantee. If the permutation null has near-zero dispersion, the placebo is treated as uninformative for that threshold rather than as calibrated evidence for or against predictability. Such a diagnostic is reported as N/A: it is not counted as a permutation failure, and it is not sufficient for a full panel-rule pass.

The factor-alpha test estimates

$$R_t = \alpha + \beta_M(MKT_t - RF_t) + \beta_S SMB_t + \beta_H HML_t + \beta_U MOM_t + \varepsilon_t, \quad (7)$$

using HAC standard errors. The null is $H_0 : \alpha \leq 0$. Including momentum is essential because the public signal is itself momentum-like. The factor-alpha regression uses frequency-matched public factor returns and the same HAC lag rule as the main Sharpe analysis. The alpha table is a diagnostic rather than the main test for canonical benchmark validation.

4.3 Turnover-scaled costs

For the equal-weighted threshold portfolio, define target weights

$$w_{t,i} = \frac{\text{sign}(s_{t,i}) \mathbf{1}\{i \in \mathcal{S}_t\}}{|\mathcal{S}_t|}, \quad w_{t,i} = 0 \text{ if } |\mathcal{S}_t| = 0. \quad (8)$$

Daily turnover is the half- L^1 distance between adjacent target weight vectors,

$$TO_t = \frac{1}{2} \sum_i |w_{t,i} - w_{t-1,i}|, \quad \widehat{TO} = T_{TO}^{-1} \sum_t TO_t, \quad (9)$$

where the average is over adjacent dates with at least one active side. For cost c per full rebalance in decimal units, the cost-adjusted Sharpe ratio is

$$\widehat{SR}_{net}(c) = \frac{\hat{\mu} - \widehat{TO} c}{\hat{\sigma}} \sqrt{A_f}. \quad (10)$$

The cost term $\widehat{TO}c$ is subtracted from the period mean return before annualization; $\widehat{TO}$ is averaged at the source frequency, daily for the public panel sources used here. This is a first-order cost stress. It assumes linear cost per unit of turnover and holds the return

volatility fixed after subtracting costs. These assumptions are deliberately simple and are used only as a screening stress. The resulting cost drag may understate implementation costs when liquidity demand is material. They are not a full execution model for high-turnover or illiquid strategies. Such applications require spread, latency, borrow, and position-size models before any implementation claim. This is a conservative cost-stress screen, not an execution model. The empirical tables also report the break-even cost in basis points, defined by $\hat{\mu} - \widehat{TOC} = 0$. The public sorted-portfolio examples do not contain average daily volume or order-size information, so they do not support an execution-scale claim.

4.4 Decision rule for reported panel candidates

The implementation reports component diagnostics, and the pre-specified summary rule is deliberately stricter than any single statistic. The panel decision rule is a conservative reporting rule: a candidate is first evaluated at the date-level statistical boundary, then with dependent-bootstrap robustness, researcher-menu adjustment, same-date placebo evidence, and turnover-cost sensitivity. A candidate is reported as passing the full decision rule only if all applicable diagnostics are informative and it clears the pre-specified thresholds:

$$p_{\text{boot}} \leq \alpha, \quad p_{\text{perm}} \leq \alpha, \quad p_{\text{HAC}} \leq \alpha, \quad p_{\text{RW}} \leq \alpha, \quad \widehat{\text{SR}}_{\text{net}}(c) > 0.$$

If the same-date permutation null is degenerate or otherwise uninformative, the candidate is reported separately as passing or failing the remaining applicable checks; it is not classified as a full decision-rule pass. This is a conservative reporting policy, not a uniformly most powerful test, an optimal weighted test, or a claim that the checks exhaust all possible failure modes. The checks are conjunctive because each one addresses a different reason a performance claim can fail: bootstrap-fragile Sharpe evidence, a placebo-consistent signal, dependence-adjusted insignificance, researcher-menu selection, or nonpositive net Sharpe after turnover-cost sensitivity. Applications with different decision costs may choose different rules ex ante, but they should not choose the rule sequence after seeing the results. The component diagnostics are reported so readers can apply less conservative decision rules. The reported tables make the threshold reproducible: in addition to the conventional $\alpha = 0.05$, a reader can recompute pass counts at stricter levels such as $\alpha = 0.01$ without changing the statistical objects being tested. The same-date permutation check is a panel placebo: it is applicable only when row-level signals can be permuted within date. Pre-aggregated factor series can still be economically interesting and can pass HAC-delta or Romano-Wolf tests, but they do not contain the row-level signal panel needed for this placebo. Such cases are reported as not applicable for the permutation check rather than as passed panel evaluations.

The stages have different interpretations. The permutation diagnostic asks whether the signal-return pairing beats a same-date placebo. The HAC diagnostic asks whether the resulting date-return Sharpe is positive after serial dependence and heteroskedasticity are accounted for. The bootstrap diagnostic asks whether that conclusion survives re-sampling of the date-return series. The Romano-Wolf diagnostic asks whether the claim survives the fixed researcher menu. The cost diagnostic reports whether net Sharpe remains positive after turnover-scaled frictions. Factor-alpha, holdout, subperiod, and stationarity diagnostics are reported alongside these checks to prevent a single favorable statistic from becoming the claim.

5 Sampling-boundary methodology

5.1 Mean-return special case

Assumption 1 (Equicorrelated row model). *For each date t , selected row returns satisfy*

$$y_{t,i} = \mu + \lambda_t + \epsilon_{t,i},$$

where λ_t is a common date shock, $\epsilon_{t,i}$ is idiosyncratic, λ_t is uncorrelated with $\epsilon_{t,i}$, $\text{Var}(\lambda_t) = \sigma_\lambda^2$, and $\text{Var}(\epsilon_{t,i}) = \sigma_\epsilon^2$. The selected count is $n_t = \bar{n}$ for the displayed derivation, and dates are independent.

Under this model,

$$\sigma^2 = \text{Var}(y_{t,i}) = \sigma_\lambda^2 + \sigma_\epsilon^2, \quad \rho = \text{Corr}(y_{t,i}, y_{t,j}) = \frac{\sigma_\lambda^2}{\sigma_\lambda^2 + \sigma_\epsilon^2} \quad (i \neq j).$$

Thus ρ is the intraclass correlation induced by the common date shock; economically, it is the fraction of row-return variance explained by shocks shared by same-date selections. In empirical diagnostics, ρ is estimated after the candidate threshold is fixed. For each panel candidate, the selected signed row returns are aligned by date and asset or portfolio, and $\hat{\rho}$ is the average off-diagonal sample correlation among active row-return series with overlapping dates; $\bar{n}$ is the average selected count across active dates. When the selected set varies over time, this equicorrelation estimate is a descriptive compression of the same-date covariance structure, not an identifying assumption for the confirmatory HAC or bootstrap tests.

Under the equicorrelated row model, the corresponding mean-return design-effect calculation gives the ratio of the naive row-pooled reported standard error of the mean to the date-level standard error:

$$\frac{SE_{\text{row,naive}}}{SE_{\text{date,true}}} = \{1 + (\bar{n} - 1)\rho\}^{-1/2}.$$

For $\rho > 0$ and $\bar{n} > 1$, the naive row-level standard-error report is too small.

This calculation is not new as a sampling result; it is the Moulton-like special case of the paper’s UVIF logic, written for signed trading-panel rows and date-level portfolio returns. To avoid conflict with the researcher menu notation $\mathcal{M}$, write the cross-sectional variance design effect as

$$\text{UVIF}_{\mu, xs}(\bar{n}, \rho) = 1 + (\bar{n} - 1)\rho.$$

The square root is reported when the object is a standard-error or t -statistic adjustment:

$$\text{UVIF}_{\mu, xs}^{1/2}(\bar{n}, \rho) = \sqrt{1 + (\bar{n} - 1)\rho}.$$

Thus UVIF denotes a variance ratio, while $\text{UVIF}^{1/2}$ denotes the associated standard-error multiplier. This mean-return calculation is expository: it makes the scale of understatement visible before the paper turns to HAC and dependent-bootstrap decisions on the realised date-return series. The magnitude can be economically large. The relevant inflation factor for the true standard error relative to the naive row report is $\text{UVIF}_{\mu, xs}^{1/2}(\bar{n}, \rho)$. With $\bar{n} = 50$, increasing same-date correlation from $\rho = 0.10$ to $\rho = 0.60$ raises this factor from 2.43 to 5.51. The inflation factor rises mechanically when selected counts and same-date common-shock correlation rise.

The proof is given in the technical appendix.

Remark 1 (Heterogeneous factor covariance). *The equicorrelated model is a closed-form illustration, not a restriction imposed by the protocol. Let a same-date signed-return vector satisfy*

$$y_t = \mu \mathbf{1} + B f_t + u_t,$$

where B contains asset-specific loadings, $\text{Var}(f_t) = \Sigma_f$, and $\text{Var}(u_t) = D_t$. Then the same-date covariance matrix is

$$\Sigma_t = B \Sigma_f B' + D_t,$$

with heterogeneous pairwise correlations whenever loadings or idiosyncratic variances differ. For any displayed portfolio weights w_t , the relevant one-date variance is $w_t' \Sigma_t w_t$, not the row-independent variance. The method avoids estimating the full Σ_t as a prerequisite for inference by moving the evidence to the realised date-return series and then using HAC or dependent bootstrap inference on that series. The design-effect calculation is therefore best read as the transparent special case that explains why row pooling fails.

5.2 Unified variance inflation for the Sharpe ratio

The mean-return design effect is useful but incomplete because the reported performance metric is usually the Sharpe ratio. The Sharpe ratio is a nonlinear functional of the first

two moments, so its variance inflation depends on the covariance of R_t and R_t^2 , not only on the variance of R_t . This motivates the Unified Variance Inflation Factor:

$$\text{UVIF}_{\text{SR}} = \frac{\text{Var}_{LR}(\widehat{\text{SR}}_{\text{portfolio}})}{\text{Var}_{\text{IID}, \text{row}}(\widehat{\text{SR}}_{\text{row}})}. \quad (11)$$

The numerator is the long-run delta-method variance at the date-return boundary. The denominator is the row-pooled delta-method variance that would be reported if the selected rows were treated as independent observations. Equation (11) is therefore a reported-uncertainty distortion ratio. It compares the uncertainty of the economically relevant date-level Sharpe claim with the uncertainty that would be reported under row-pooled IID analysis. Under exchangeability of the linearized Sharpe influence function, and when the row-pooled and date-level estimators target the same equal-weight Sharpe functional, the cross-sectional part of this variance ratio reduces to the same $1 + (\bar{n} - 1)\rho$ design effect shown above, with ρ applied to the linearized Sharpe influence rather than to raw returns. Outside that special case, Sharpe UVIF is target-aware: it measures the gap between the row-level report and the date-level claim that can be interpreted as portfolio-performance evidence.

Assumption 2 (Sharpe UVIF regularity). *The date-level process $G_t = (R_t, R_t^2)'$ is stationary with finite $4 + \delta$ moments, $v_R = m_{2,R} - \mu_R^2 \geq \eta > 0$, and finite long-run covariance Ω_R . The aggregation rule is fixed before evaluation. The displayed decomposition uses a constant selected-count approximation $\bar{n}$, and the row-pooled IID denominator is a counterfactual reporting benchmark rather than the maintained sampling law.*

Proposition 1 (Sharpe UVIF decomposition). *Let R_t be the date-level portfolio excess return and $G_t = (R_t, R_t^2)'$. Define $\theta_R = (\mu_R, m_{2,R})' = \mathbb{E}[G_t]$, $v_R = m_{2,R} - \mu_R^2 > 0$. For source-specific annualization factor A_f , define*

$$d_R = \nabla f(\theta_R) = \left(\sqrt{A_f} \frac{m_{2,R}}{v_R^{3/2}}, -\frac{1}{2} \sqrt{A_f} \frac{\mu_R}{v_R^{3/2}} \right)'.$$

Let $\Gamma_k^R = \mathbb{Cov}(G_t, G_{t-k})$ and $\Omega_R = \Gamma_0^R + \sum_{k \geq 1} (\Gamma_k^R + \Gamma_k^{R'})$. For the selected row return $Y_{t,i}$, let $H_{t,i} = (Y_{t,i}, Y_{t,i}^2)'$, $\theta_Y = \mathbb{E}[H_{t,i}]$, $d_Y = \nabla f(\theta_Y)$, and $\Psi_Y = \text{Var}(H_{t,i})$. With a constant selected count $\bar{n}$, the first-order Sharpe variance inflation relative to a row-pooled IID report is

$$\text{UVIF}_{\text{SR}} = \underbrace{\frac{\bar{n} d_R' \Gamma_0^R d_R}{d_Y' \Psi_Y d_Y}}_{\text{cross-sectional UVIF component}} \times \underbrace{\frac{d_R' \Omega_R d_R}{d_R' \Gamma_0^R d_R}}_{\text{Sharpe long-run serial factor}}. \quad (12)$$

Moreover, the serial factor can be written as

$$1 + 2 \sum_{k=1}^{\infty} \psi_k^{\text{SR}}, \quad \psi_k^{\text{SR}} = \frac{d'_R \Gamma_k^R d_R}{d'_R \Gamma_0^R d_R},$$

and the scalar long-run variance entering the numerator is

$$d'_R \Omega_R d_R = \frac{A_f}{v_R^3} \left[m_{2,R}^2 \Omega_{11} - \frac{\mu_R m_{2,R}}{2} (\Omega_{12} + \Omega_{21}) + \frac{\mu_R^2}{4} \Omega_{22} \right], \quad (13)$$

where Ω_{ab} denotes the (a, b) element of Ω_R .

The proof is given in the technical appendix.

5.3 What Proposition 1 does and does not identify

Proposition 1 is a first-order variance decomposition for a fixed aggregation rule and a constant selected-count approximation. It identifies how a row-pooled IID Sharpe uncertainty report changes when the claim is moved to the date-return boundary and when the date series has long-run serial covariance in (R_t, R_t^2) . It does not identify a causal trading edge, does not prove that row and date Sharpe functionals are equal outside the stated exchangeability conditions, and does not replace the confirmatory date-level HAC, dependent-bootstrap, researcher-menu, placebo, factor-alpha, and cost diagnostics. The empirical role of Sharpe UVIF is therefore diagnostic: it explains the sampling-boundary distortion that the reported diagnostics then evaluate with date-level inferential procedures. The sample UVIF reported in the empirical section is a plug-in diagnostic for this population ratio; confirmatory rejection decisions do not depend on estimating UVIF consistently.

Remark 2 (Connection to Moulton-like design effects). *If the linearized row Sharpe influence $\zeta_{t,i}^Y = d'_Y(H_{t,i} - \theta_Y)$ is exchangeable within date with intraclass correlation ρ_ζ , and the date-level linearization is the equal-weighted average of those row influences, the cross-sectional factor in (12) reduces to $1 + (\bar{n} - 1)\rho_\zeta$. The mean-return design effect is the special case in which the influence variable is $Y_{t,i} - \mu$. The Sharpe case replaces raw return correlation by correlation of the linearized nonlinear-performance influence, which is why skewness, kurtosis, and volatility clustering enter the variance inflation.*

Fixed-menu validity. The HAC-delta and block-bootstrap procedures are used under the standard conditions for a fixed finite researcher menu: stationary date-return series, adequate moments for the Sharpe delta-method expansion, non-degenerate volatility, and block lengths satisfying $\ell \rightarrow \infty$ with $\ell/T \rightarrow 0$. Under those conditions, standard delta-method and block-bootstrap arguments support marginal Sharpe inference for each

fixed menu element, and Holm or Romano-Wolf stepdown procedures control familywise error for the fixed menu (Holm, 1979; Kunsch, 1989; Romano and Wolf, 2005; Politis and Romano, 1994; Lahiri, 2003). These standard conditions do not cover long memory, infinite-variance returns, growing menus, structural breaks, or nonstationary date returns. The empirical stationarity, holdout, placebo, factor-alpha, and cost outputs are therefore diagnostics rather than proof that every assumption holds in every market setting.

6 Synthetic positive-control calibration

The reporting rule is not mechanically anti-discovery. To check that the decision criteria can pass a true signal, a synthetic panel is generated with a persistent predictive signal, moderate same-date dependence, and serial dependence in the date-level returns. The signal is constructed before inference is run, the same date-aggregation and reporting criteria are applied, and the resulting positive-control candidate passes HAC, Romano-Wolf, and same-date permutation checks. This calibration is not used as evidence for a public trading strategy; it shows that the procedure can accept known positive evidence when the data-generating process contains a genuine date-level edge.

The positive-control DGP uses $T = 5000$ dates and $N = 100$ entities, annual alpha 0.12 (daily alpha 0.12/252), idiosyncratic volatility 0.01, common-factor loading 0.50, AR(1) common-factor coefficient 0.10, signal-stay probability 0.98, and seed 777. The generated output records the bootstrap and permutation counts used for the detailed run and reports a lower-resolution replication screen as a power sanity check. The design is intentionally simple: it is a representative known-edge calibration check for the reporting criteria, not a model of any public anomaly.

Table 1: Synthetic positive-control evaluation.

Panel	$\bar{\rho}$	SE infl.	SR	HAC p^+	RW p	Perm. p^+
Synthetic control	0.202	4.58	1.744	7.67e-13	2.00e-04	1.00e-04

Table 2: Positive-control replication screen.

Reps	Boot	Perms	Block pass	HAC pass	RW pass	Perm. pass	Cost pass	All pass
100	499	199	100.0%	100.0%	100.0%	100.0%	100.0%	100.0%

Note: The replication screen uses fewer resampling draws than the detailed calibration row and is reported only as a power sanity check.

Appendix material. The worked numerical example, full proof details, positive-control data generation notes, and secondary robustness tables are reported in the technical appendix. The main manuscript keeps only the decision rule and the evidence needed to evaluate it.

7 Empirical results: benchmark illustrations

The public evidence is illustrative rather than a broad anomaly-panel survey. It evaluates the canonical Kenneth French winner-minus-loser momentum decile benchmark, a broader Kenneth French decile-rank threshold profile, Kenneth French 25 Size/BM portfolios with a skipped-momentum signal, and public AQR value, momentum, quality, beta, and HML-style factor candidates. The AQR series are economically meaningful factor-series benchmarks, but they are pre-aggregated series rather than row-level constituent panels; when the same-date permutation check cannot be applied, the benchmark is outside the full panel rule. The main tables therefore separate panel candidates from single-series factor benchmarks. Failed public-data access attempts are kept in the provenance files rather than treated as main empirical results.

Table 3 reports row-level panel candidates with the structure needed for the full panel decision rule. Table 4 reports pre-aggregated factor benchmarks subject to time-series inference only. Table 5 verifies that the canonical WML candidate matches the direct French winner-minus-loser decile benchmark. Table 6 compares the row-naive report, date-level HAC inference, dependent-bootstrap inference, and Sharpe UVIF for the panel candidates. Table 7 gives the sampling-boundary summary using row p -values, displayed UVIF, conservative maximum p -values, and net Sharpe status for panel candidates. Table 8 reports the targeted $L = 1, 5, 21$ horizon-effect sensitivity for the dynamic Size/BM momentum panel. The technical appendix reports the full secondary table set: sampling-boundary diagnostics, same-date permutation details, HAC bandwidth/prewhitening/fixed- b robustness, rule sensitivity, holdout splits, and cost sweeps. The reproducibility package records the fixed public sources, candidate types, source frequencies, annualization factors, thresholds, lookbacks, skips, eligibility rules, and configuration hashes used in the reported tables.

The factor-alpha, data-snooping, stationarity, holdout, and cost tables are secondary diagnostics, not replacements for the primary date-return decision. A raw positive Sharpe that does not survive FF3 plus momentum alpha, bootstrap robustness, Romano-Wolf/White checks, same-date placebos, or plausible costs should not be interpreted as reliable portfolio-performance evidence. This is the empirical implication of the sampling-boundary design: familiar public candidates can look attractive at one boundary and fail at the date-level decision boundary.

The panel-candidate table reports only sources with row-level signal-return structure and therefore eligibility for the full panel decision rule. The single-series benchmark table reports pre-aggregated factor returns separately; those rows are time-series benchmarks, not failed panel candidates.

The benchmark validation table shows that the panel implementation matches the direct French winner-minus-loser construction.

Table 3: Panel candidates with row-level signal-return structure.

Panel candidate	SR	HAC p+	Boot p+	RW p	Perm p	Net SR	Full rule 5%	Applicable checks 5%
French momentum deciles, rank-threshold panel	0.447	7.54e-05	2.00e-04	2.00e-04	N/A	0.447	no	yes
French WML decile benchmark	0.497	5.53e-06	2.00e-04	2.00e-04	1.00e-04	0.497	yes	yes
French 25 Size/BM, skipped-momentum panel	0.058	0.412	0.422	0.480	0.943	-0.050	no	no

Table 4: Single-series factor benchmarks subject to time-series inference only.

Single-series benchmark	SR	HAC p+	Boot p+	RW p	Scope
AQR Betting Against Beta, equity factors	0.741	3.08e-13	2.00e-04	2.00e-04	time-series only
AQR Quality Minus Junk	0.536	2.26e-05	2.00e-04	2.00e-04	time-series only
AQR HML Devil	0.336	0.004	0.002	0.002	time-series only
AQR Value and Momentum Everywhere	0.572	1.62e-04	2.00e-04	2.00e-04	time-series only

Note: Time-series only means not eligible for same-date permutation, row-retention diagnostics, or the full panel decision rule.

Table 5: Canonical French momentum benchmark validation.

Candidate	thr	n	start	end	panel SR	direct WML SR	scale
French WML decile benchmark	0.900	26131	1926-11-04	2026-04-30	0.497	0.497	0.500

The comparator table makes the standard-practice contrast explicit: row-naive evidence is shown next to date-level HAC and bootstrap inference on the portfolio return series.

Table 6: Comparison with common row and date-level inference choices for panel candidates.

Panel candidate	row p+	date-IID p+	date-HAC p+	block p+	UVIF
French momentum deciles, rank-threshold panel	1.70e-04	2.00e-04	7.54e-05	2.00e-04	1.000
French WML decile benchmark	4.85e-04	2.00e-04	5.53e-06	2.00e-04	1.000
French 25 Size/BM, skipped-momentum panel	0.253	0.417	0.412	0.422	9.002

Note: The row-naive column treats selected rows as independent. Date-IID, date-HAC, and block-bootstrap columns operate on the date-level portfolio return series.

The sampling-boundary table summarizes the boundary directly: row-level evidence can disappear once same-date redundancy, placebo evidence, and turnover-scaled net Sharpe are evaluated together.

Table 7: Sampling-boundary summary for panel candidates.

Candidate	Row p+	UVIF	Max p+	Net SR	Status
French momentum deciles, rank-threshold panel	1.70e-04	1.000	2.00e-04	0.447	Permutation uninformative; applicable checks pass
French WML decile benchmark	4.85e-04	1.000	2.00e-04	0.497	Passes full rule
French 25 Size/BM, skipped-momentum panel	0.253	9.002	0.943	-0.050	Fails statistical and cost checks

Note: Displayed UVIF is a one-sided inflation summary floored at one. Degenerate same-date permutation diagnostics are reported as uninformative rather than converted into failures.

The horizon table shows that changing the lookback reduces turnover in the Size/BM stress panel but does not convert it into a robust net-Sharpe claim.

The public illustrations separate benchmark validation from broader search claims. The canonical French WML candidate is expected to match the direct Kenneth French

Table 8: Horizon-effect sensitivity for the dynamic Size/BM momentum panel.

L	thr	dates	avg names	gross SR	net SR 5bps	turnover	UVIF SR	serial UVIF	HAC p+
1	0.500	3497	12.913	0.073	-0.394	0.733	11.498	0.891	0.386
5	0.500	3493	13.000	-0.312	-0.528	0.339	12.880	0.990	0.878
21	0.500	3477	13.000	0.058	-0.050	0.169	12.121	0.933	0.412

winner-minus-loser decile construction and is reported as a public benchmark check. Its full-rule pass does not undermine the sampling-boundary argument; it shows that date-level inference can preserve canonical evidence when the signal is strong enough and the placebo, menu, dependence, and cost diagnostics do not overturn it. In the reported output, the WML panel Sharpe matches the direct WML Sharpe at approximately 0.50; the scale difference is exactly one-half because the panel portfolio holds equal long and short legs while the direct WML series is winner minus loser. The broader rank-threshold profile and the French Size/BM dynamic-momentum panel are evaluated as stress cases rather than as canonical anomaly replications. The AQR factor-series candidates have positive time-series evidence in the reported HAC and block-bootstrap diagnostics, but because they are pre-aggregated factor returns rather than row-level signal panels, they are not classified as passed panel-evaluation candidates.

The horizon-effect sensitivity shows that the dynamic Size/BM momentum stress case fails because of weak signal and execution drag, not because the canonical momentum benchmark is missing. At threshold 0.5, $L = 1$ has very high turnover and a negative 5 bps net Sharpe; $L = 21$ lowers turnover substantially but still leaves net Sharpe slightly negative. The empirical Sharpe UVIF relative to row-naïve inference is about 11–13 across the three horizons, while the serial-only UVIF is close to one in this sample. Thus the dominant inflation channel in this public stress panel is cross-sectional row multiplication; a longer lookback does not mechanically make serial dependence the dominant force. The sampling-boundary summary is intentionally limited to rows with the needed public panel structure. It does not assert that every named factor is a panel candidate. Pre-aggregated AQR factor series are excluded from that table because they do not provide row-level signal-return pairs from which row p -values, same-date permutation tests, or panel row-retention diagnostics can be computed.

8 Monte Carlo evidence: 35 percent row-naïve rejection

The Monte Carlo study is designed as a size-and-power check of the sampling boundary, not as a search for a favorable strategy. Each replication generates a symbol-date return panel, applies the same date-aggregation rule used in the empirical analysis, and then

asks whether each inference method rejects a positive-edge null at the nominal 5 percent level. The base panel is

$$r_{t,i} = \mu_t + \beta_i' f_t + \epsilon_{t,i}, \quad f_t = \Phi f_{t-1} + u_t.$$

The grid varies serial dependence, cross-sectional dependence, GARCH volatility, multi-factor exposure, and regime switching. The high-persistence AR(1) case is reported as such; it is not labeled as a long-memory design. The high-volatility GARCH case is not described as an infinite-variance heavy-tail design. The simulation uses 1000 Monte Carlo replications for the displayed size and rejection-rate tables, $T = 1000$ dates and $N = 50$ entities in the main null designs, 5000 bootstrap draws per simulation, and seed-indexed replications generated from the simulation index. Moving-block and stationary bootstrap procedures use the block-length rule reported in Table 10: for moving-block bootstrap, $\ell = \lceil T^{1/5} \max(D, 1) \rceil$, capped at $\lfloor T/3 \rfloor$, and the stationary bootstrap uses the same mean block length. The full DGP grid, seed policy, configuration hash, and run metadata are included in the reproducibility package. The null designs set the true edge to zero while varying cross-sectional correlation and serial memory; the power designs then increase the true annualized Sharpe on the same dependence-aware decision boundary. The important object is the gap between the row-naïve rejection rate and the rejection rates of methods that operate on date-level returns. In the displayed null grid, the row-naïve test overrejects around 35–37 percent under dependent nulls, while HAC-delta, moving-block bootstrap, stationary bootstrap, and Romano-Wolf remain much closer to nominal size. In the power grid, row-naïve rejection reaches about 50 percent at a true annualized Sharpe near 0.19, before the dependence-aware methods show comparable power; that gap reflects null overrejection rather than valid incremental power.

The size and power experiment compares row-level naïve testing, date-IID bootstrap, HAC-delta inference, moving-block bootstrap, stationary bootstrap, and Romano-Wolf joint resampling. The coverage experiment focuses on the date-level Sharpe target, so row-naïve intervals are interpreted as coverage of the economically relevant date-level claim rather than coverage of a separate row-distribution Sharpe. The purpose is not to rank methods by power alone. A method that appears powerful because it overrejects under the null is not acceptable for confirmatory evaluation.

Table 9 states the main design contract for the null and positive-control experiments, and Table 10 reports the resampling settings used for the reported results. Table 11 gives the main size evidence. The row-naïve column is the familiar symbol-date analysis; the HAC-delta, moving-block, stationary, and Romano-Wolf columns make decisions from the date-return series. When row-naïve rejection is high in null or near-null designs but the dependence-aware methods remain near nominal size, the apparent extra power is a false precision gain from treating same-date rows as independent evidence. Table

12 reports row-naive interval behavior against the date-level Sharpe target, making the target mismatch explicit. Table 13 then sweeps same-date correlation and selected-count pressure under a zero-edge null. These sweeps are diagnostic stress tests for the sampling boundary; they do not introduce a new empirical claim. The technical appendix reports the full DGP grid, coverage table, power table, and power-curve figure.

The null-size figure and table are the core simulation results: under dependent nulls, row-naive testing rejects about one third of the time at a nominal 5 percent level, while date-level HAC, block-bootstrap, stationary, and Romano-Wolf procedures stay much closer to size.

Table 9: Monte Carlo design contract.

Design	T	N	Selected count	XS dep.	Serial dep.	Heterosk.	Reps	Resampling	Seed
IID null	1000	50	all rows	rho=0	AR(1)=0	none	1000	5000 boot	rep index
Dependent null	1000	50	all rows	rho=0.2	AR(1)=0.2	none	1000	5000 boot	rep index
GARCH null	1000	50	all rows	rho=0.2	AR(1)=0.2	GARCH(0.05,0.90)	1000	5000 boot	rep index
Positive control	5000	100	signal-selected rows	common loading=0.50	AR(1)=0.10	none	100	499 boot	777

Note: All null designs have true Sharpe equal to zero. Positive-control rows use a declared nonzero signal and are included to show that the reporting rule is not mechanically anti-discovery.

Table 10: Monte Carlo and resampling settings.

Setting	Value
Monte Carlo replications	1000
Bootstrap draws	5000
Permutation draws	10000
Main null panel sizes	T=1000, N=50
Cross-sectional correlation grid	0.000, 0.200, 0.250, 0.300, 0.350, 0.400, 0.500
AR(1) grid	0.000, 0.150, 0.200, 0.300, 0.500, 0.700
GARCH parameter pairs	(0.050, 0.900); (0.100, 0.850); (0.150, 0.800)
HAC bandwidth rule	automatic Bartlett, with fixed-lag sensitivity checks
Moving-block length	ceil(T to the 1/5 times max(D,1)), capped at floor(T/3)
Stationary mean block	same length as moving-block rule
Dependence factor D	(2 rho1/(1-rho1)) to the 2/3 for positive AR(1) rho1; otherwise 1
Seed policy	seed-indexed replications from campaign runner

Table 11: Null rejection rates at nominal 5 percent size.

dgp	date_iid	hac_delta	moving_block	romano_wolf	row_naive	stationary
null_dep	0.094	0.066	0.070	0.068	0.352	0.067
null_garch	0.098	0.063	0.070	0.067	0.351	0.068
null_iid	0.060	0.060	0.061	0.060	0.062	0.062

9 Discussion and limitations

The results should be read as a financial-econometrics testing exercise for entity-time panel evidence. The framework addresses two common sources of false precision: same-




Figure 2: Rejection rates under null designs. The dashed line is the nominal 5 percent rejection rate. Row-naive rejection is near nominal under IID sampling but rises sharply under dependent nulls, while date-level dependence-aware procedures remain much closer to nominal size.

Table 12: Row-naive interval behavior when evaluated against the date-level Sharpe target.

DGP	Name	row coverage	bias vs date SR	CI width	mean SE	n
01_iid	IID baseline	0.000	-2.567	0.278	0.071	1000
02_ser_mild	Serial mild	0.000	-2.567	0.278	0.071	1000
03_ser_strong	Serial strong	0.000	-2.568	0.278	0.071	1000
04_xs_mild	Cross mild	0.073	-0.474	0.278	0.071	1000
05_xs_strong	Cross strong	0.262	-0.153	0.278	0.071	1000
06_both_mod	Both moderate	0.088	-0.503	0.278	0.071	1000
07_both_strong	Both strong	0.178	-0.145	0.278	0.071	1000
08_realistic	Realistic	0.107	-0.277	0.124	0.032	1000
09_garch_mild	GARCH mild	0.110	-0.473	0.278	0.071	1000
10_garch_strong	GARCH strong	0.199	-0.268	0.278	0.071	1000
11_multifactor	Multi-factor K=3	0.225	-0.327	0.278	0.071	1000
12_regime	Regime-switching	0.306	-0.092	0.197	0.050	1000
13_high_ar1	High-persistence AR(1)	0.137	-0.338	0.278	0.071	1000
14_garch_highvol	High-volatility GARCH	0.193	-0.325	0.278	0.071	1000

Note: These are not ordinary row-estimator coverage rates; they evaluate the row-naive interval against the economically relevant date-level target.

Table 13: Sampling-boundary stress tests under a zero-edge null.

Panel	Setting	Row rej.	HAC rej.	Row mean SE	HAC mean SE
A. Same-date correlation (N=50)	rho=0.000	0.097	0.056	0.071	0.598
A. Same-date correlation (N=50)	rho=0.100	0.310	0.064	0.071	0.598
A. Same-date correlation (N=50)	rho=0.300	0.354	0.044	0.071	0.599
A. Same-date correlation (N=50)	rho=0.350	0.368	0.071	0.071	0.597
A. Same-date correlation (N=50)	rho=0.500	0.409	0.047	0.071	0.597
A. Same-date correlation (N=50)	rho=0.700	0.408	0.060	0.071	0.599
B. Selected count (rho=0.35)	N=5	0.192	0.058	0.225	0.598
B. Selected count (rho=0.35)	N=25	0.308	0.052	0.100	0.600
B. Selected count (rho=0.35)	N=50	0.368	0.071	0.071	0.597
B. Selected count (rho=0.35)	N=100	0.421	0.056	0.050	0.598

Note: Row-naive overrejection rises with same-date dependence and selected-count pressure, whereas date-level HAC remains close to nominal size.

date cross-sectional row multiplication and serial dependence in the date-level return series. It does not remove the need for economic interpretation, out-of-sample validation, liquidity analysis, or venue-specific execution modeling.

The formal validity statement assumes a fixed and finite researcher menu, stationary date-return series, adequate moments for the Sharpe delta-method expansion, and block-bootstrap regularity. The framework does not cover long-memory processes, infinite-variance returns, growing researcher menus, structural breaks, or nonstationary date returns without additional assumptions. The empirical stationarity, holdout, factor-alpha, placebo, and cost outputs are therefore diagnostics and boundary checks, not proofs that every maintained assumption holds in every market setting. A locked holdout period or independently refreshed sample after the researcher menu is frozen remains the stronger design when such data are available.

10 Data and code availability

The reproducibility package contains the LaTeX source, generated public aggregate tables, public-data loaders, simulation code, unit tests, release scripts, and SHA256 provenance. The public repository is: Repository link omitted from the review manuscript to preserve double-blind anonymity; it is supplied in the separate title page and submission metadata.

Code is released under Apache-2.0 and manuscript/public aggregate tables under CC-BY 4.0 unless a target venue requires different terms. Raw third-party data are not redistributed where source terms do not permit redistribution; the repository instead provides documented loaders, source links, access metadata, generated aggregate tables, and checksums needed to reproduce the analysis. Proprietary predictions, raw trades, private tickers, production feature definitions, and unrelated workspace files are excluded. The technical appendix contains the full simulation grid, robustness tables, positive-control output, configuration notes, and proof details omitted from the main manuscript for length.

11 Conclusion

The paper’s central message is simple. In financial panels, multiplying same-date rows does not automatically create independent Sharpe-ratio evidence when the relevant payoff is a date-level portfolio return. The solution proposed here is not a new anomaly but a sampling-boundary method: aggregate to the date-return series, estimate long-run Sharpe uncertainty at that boundary, and record how far the row-pooled report differed from the confirmatory date-level claim.

In public benchmark evidence, this framework preserves the canonical French WML result while overturning weaker stress-panel claims once dependence, permutation validity, researcher menus, and turnover costs are respected. Pre-aggregated AQR factor series remain useful time-series benchmarks, but without row-level constituent signals they cannot be evaluated by the panel permutation check.

The broader implication is that Sharpe-ratio evidence in panels should be judged at the date-level portfolio return boundary. Future work can extend the framework to richer implementation-cost models, time-varying cross-sectional dependence, variable selected-count asymptotics, larger anomaly panels, and richer blocked-permutation designs; those extensions are beyond the scope of this Modern Finance-focused submission.

Author contributions

The author was responsible for conceptualization, methodology, software, validation, formal analysis, data curation, writing, and revision.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Acknowledgments

Acknowledgments are omitted from the review manuscript to preserve double-blind anonymity.

Declaration of interest

The author declares no conflict of interest.

Use of AI tools

No AI tool is listed as an author. Any venue-required tool-use disclosure is supplied in the separate submission metadata.

Ethics statement

This study uses publicly available aggregate financial data and simulation data. It does not involve human participants, personal data, clinical data, animal subjects, or cell lines. Ethics approval was therefore not required.

References

- Almgren, R. and Chriss, N. (2001). Optimal execution of portfolio transactions. *Journal of Risk*, 3:5–39.
- Andrews, D. W. K. (1991). Heteroskedasticity and autocorrelation consistent covariance matrix estimation. *Econometrica*, 59(3):817–858.
- Andrews, D. W. K. and Monahan, J. C. (1992). An improved heteroskedasticity and autocorrelation consistent covariance matrix estimator. *Econometrica*, 60(4):953–966.
- Bailey, D. H. and Lopez de Prado, M. (2014). The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107.
- Bailey, D. H., Borwein, J. M., Lopez de Prado, M., and Zhu, Q. J. (2014). The probability of backtest overfitting. *Journal of Computational Finance*, 20(4):39–70.
- Barras, L., Scaillet, O., and Wermers, R. (2010). False discoveries in mutual fund performance: Measuring luck in estimated alphas. *Journal of Finance*, 65(1):179–216.
- Benjamini, Y. and Hochberg, Y. (1995). Controlling the false discovery rate. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300.
- Benjamini, Y. and Yekutieli, D. (2001). The control of the false discovery rate in multiple testing under dependency. *Annals of Statistics*, 29(4):1165–1188.
- Bouchaud, J.-P., Farmer, J. D., and Lillo, F. (2008). How markets slowly digest changes in supply and demand. In *Handbook of Financial Markets: Dynamics and Evolution*, pages 57–160.

- 648 Carhart, M. M. (1997). On persistence in mutual fund performance. *Journal of Finance*,
649 52(1):57–82.
- 650 Fama, E. F. and French, K. R. (1993). Common risk factors in the returns on stocks and
651 bonds. *Journal of Financial Economics*, 33(1):3–56.
- 652 Fama, E. F. and French, K. R. (2010). Luck versus skill in the cross-section of mutual
653 fund returns. *Journal of Finance*, 65(5):1915–1947.
- 654 Frazzini, A., Israel, R., and Moskowitz, T. J. (2015). Trading costs of asset pricing
655 anomalies. Working paper.
- 656 Harvey, C. R. and Liu, Y. (2020). False (and missed) discoveries in financial economics.
657 *Journal of Finance*, 75(5):2503–2553.
- 658 Harvey, C. R. and Liu, Y. (2020). A census of the factor zoo. SSRN working paper,
659 revised October 2020.
- 660 Harvey, C. R., Liu, Y., and Zhu, H. (2016). ... and the cross-section of expected returns.
661 *Review of Financial Studies*, 29(1):5–68.
- 662 Hasbrouck, J. (2009). Trading costs and returns for U.S. equities. *Journal of Finance*,
663 64(3):1445–1477.
- 664 Holm, S. (1979). A simple sequentially rejective multiple test procedure. *Scandinavian*
665 *Journal of Statistics*, 6(2):65–70.
- 666 Kish, L. (1965). *Survey Sampling*. Wiley.
- 667 Kiefer, N. M. and Vogelsang, T. J. (2005). A new asymptotic theory for
668 heteroskedasticity-autocorrelation robust tests. *Journal of Econometrics*, 129(1–
669 2):131–162.
- 670 Kunsch, H. R. (1989). The jackknife and the bootstrap for general stationary observations.
671 *Annals of Statistics*, 17(3):1217–1241.
- 672 Lahiri, S. N. (2003). *Resampling Methods for Dependent Data*. Springer.
- 673 Ledoit, O. and Wolf, M. (2008). Robust performance hypothesis testing with the Sharpe
674 ratio. *Journal of Empirical Finance*, 15(5):850–859.
- 675 Lo, A. W. (2002). The statistics of Sharpe ratios. *Financial Analysts Journal*, 58(4):36–
676 52.
- 677 McLean, R. D. and Pontiff, J. (2016). Does academic research destroy stock return
678 predictability? *Journal of Finance*, 71(1):5–32.

- 679 Moulton, B. R. (1986). Random group effects and the precision of regression estimates.
680 *Journal of Econometrics*, 32(3):385–397.
- 681 Novy-Marx, R. (2014). Predicting anomaly performance with politics, the weather, global
682 warming, sunspots, and the stars. *Journal of Financial Economics*, 112(2):137–146.
- 683 Opdyke, J. D. (2007). Comparing Sharpe ratios: So where are the p-values? *Journal of*
684 *Asset Management*, 8(5):308–336.
- 685 Politis, D. N. and Romano, J. P. (1994). The stationary bootstrap. *Journal of the*
686 *American Statistical Association*, 89(428):1303–1313.
- 687 Politis, D. N. and White, H. (2004). Automatic block-length selection for the dependent
688 bootstrap. *Econometric Reviews*, 23(1):53–70.
- 689 Petersen, M. A. (2009). Estimating standard errors in finance panel data sets: Comparing
690 approaches. *Review of Financial Studies*, 22(1):435–480.
- 691 Romano, J. P. and Wolf, M. (2005). Stepwise multiple testing as formalized data snooping.
692 *Econometrica*, 73(4):1237–1282.
- 693 Storey, J. D. (2002). A direct approach to false discovery rates. *Journal of the Royal*
694 *Statistical Society: Series B*, 64(3):479–498.
- 695 White, H. (2000). A reality check for data snooping. *Econometrica*, 68(5):1097–1126.